RICE UNIVERSITY

# Precision and the Anatomy of a Matrix Multiply

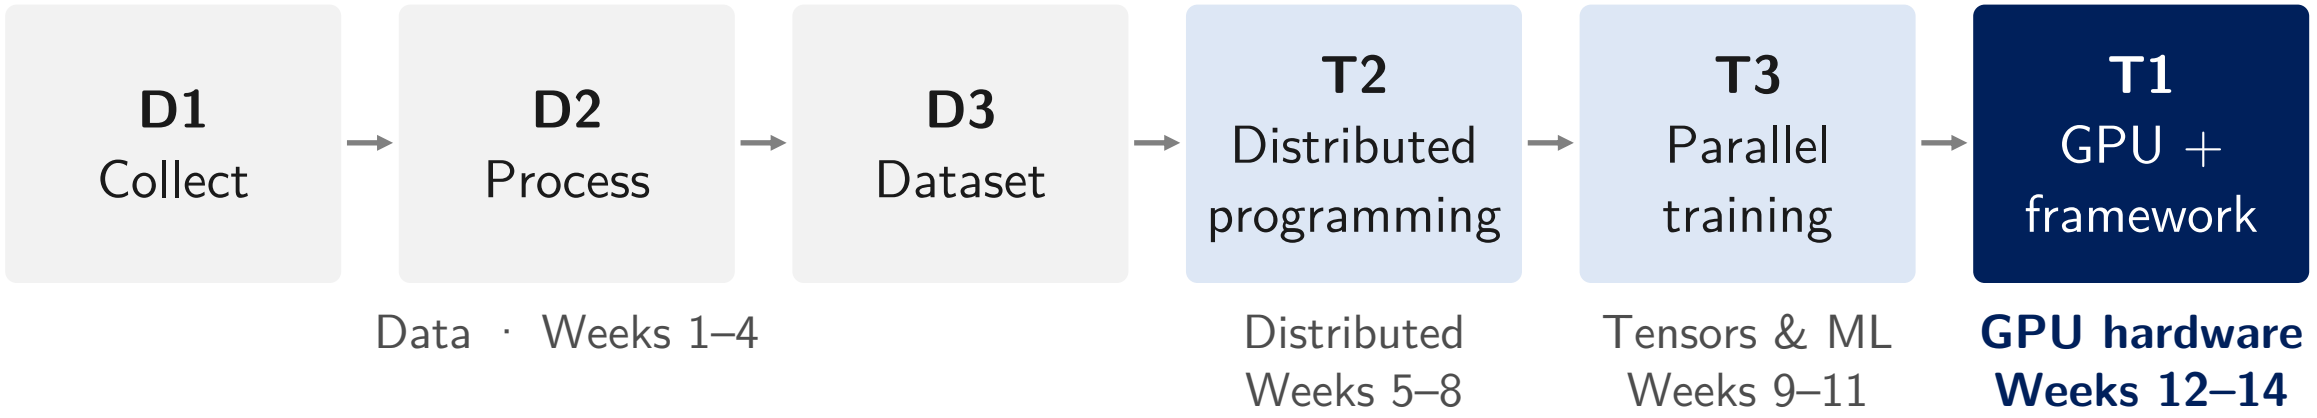
How many bits a number needs, and how a fast matrix multiply moves them

Yuke Wang

Rice University

COMP 329 · Spring 2027 · Week 14, Lecture 1

# Where we are




**In our pipeline**

Week 13 taught us to **predict** a kernel. Today: the matrix multiply — bits per number and tiles. **Lab 3** is due after today; **T1** on Friday.

# Why do we need fewer bits and tiles?

**The job:** $C = A \cdot B$, $4096 \times 4096$ matrices on one H100
$\approx$ **$1.4 \times 10^{11}$** FLOPs, almost all multiply-adds

**One thread per C[i, j], FP32, straight from HBM: how long?**

**Simplest kernel**, FP32 from HBM
$\approx$ **164 ms**

**BF16 tensor-core peak**
$\approx$ **0.14 ms**

**Key idea**

A gap of $\approx$ **1,200×**. Fewer **bits** per number and **reuse** of each number on chip close it.

Source: H100 datasheet (SXM5): 3.35 TB/s HBM, 989 TFLOP/s BF16 dense; derived values.

# Precision and tiling, in plain words

## Idea 1

**Precision**: bits per number.

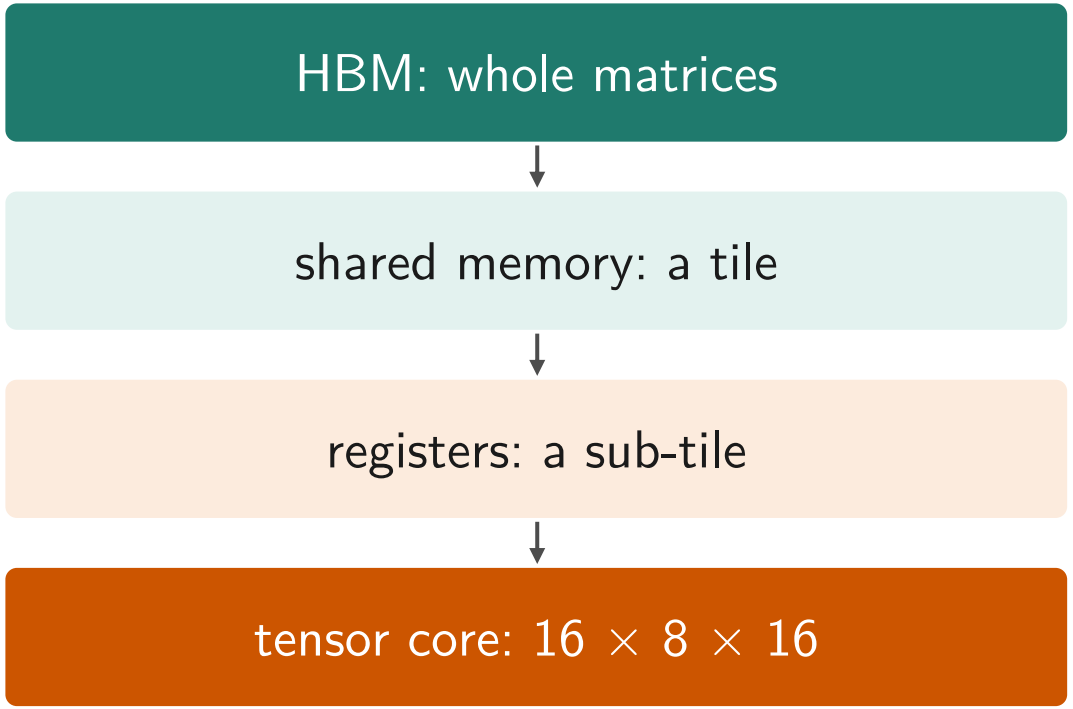
Fewer bits $\rightarrow$ fewer bytes, faster math.

Cost: less **range** or **precision**.

## Idea 2

**Tiling**: load a block on chip **once**,

then reuse it **many** times.




## Key idea

Both ideas cut **bytes moved per FLOP**, the lever the roofline (Week 13) points to.

# Where they show up

- **BF16**: the usual training format; PyTorch **autocast** picks it per operation.
- **FP8**: doubles the H100 peak again; used to train large models (DeepSeek-V3).
- **Tiling**: inside cuBLAS (behind `torch.matmul`), CUTLASS, and Triton's `tl.dot`.

**The price**

Low precision needs a **recipe** (Part 2)

Big tiles cost **shared memory** and **registers**

**In our pipeline**

**Why?** A matrix multiply is most of a training step; bits and tiles decide whether it runs near peak. Lab 3 and T1 ask you to predict exactly this.

Source: H100 datasheet (SXM5); DeepSeek-AI, "DeepSeek-V3 Technical Report," 2024; NVIDIA cuBLAS and CUTLASS docs.

# Learning goals for today

- Read a format's bits as **range** (exponent) and **precision** (mantissa): FP32, BF16, FP16, FP8.
- Explain **mixed precision**: low-precision inputs, **FP32 accumulation**, loss scaling.
- Show that a **naive** matrix multiply is **memory-bound**, and compute what **tiling** buys.
- Describe tiling through **shared memory**, **registers**, and **tensor cores**, and why tile shapes matter.

# A puzzle to start

```
s = bf16(0.0)
for _ in range(4096):
    s = bf16(s + 1.0)     # round every time
print(s)
```

**BF16**: 16 bits per number
the usual training format

## What does it print?

**Watch out**

It prints **256**. Above 256, BF16 numbers are **2 apart**, so 256 + 1 rounds back to 256.
Long sums need **more bits** than the inputs.

# Today's plan

1

## Number formats

What does a 16-bit or 8-bit number give up?

2

## Mixed precision

How do we train in 16 bits and still get FP32 answers?

3

## The anatomy of a matrix multiply

Why does a naive matmul leave the GPU idle, and how do tiles fix it?

Build

## Formats and tiles

Work in pairs, on paper. About ten minutes.

Part 1

# Number formats

What does a 16-bit or 8-bit number give up?


**1** · Number formats

**2** · Mixed precision

**3** · The anatomy of a matrix multiply

**Build** · Formats and tiles

# Anatomy of a floating-point number

A number is stored as **sign** $\times$ **(1 +** **fraction)** $\times$ **a power of 2**.

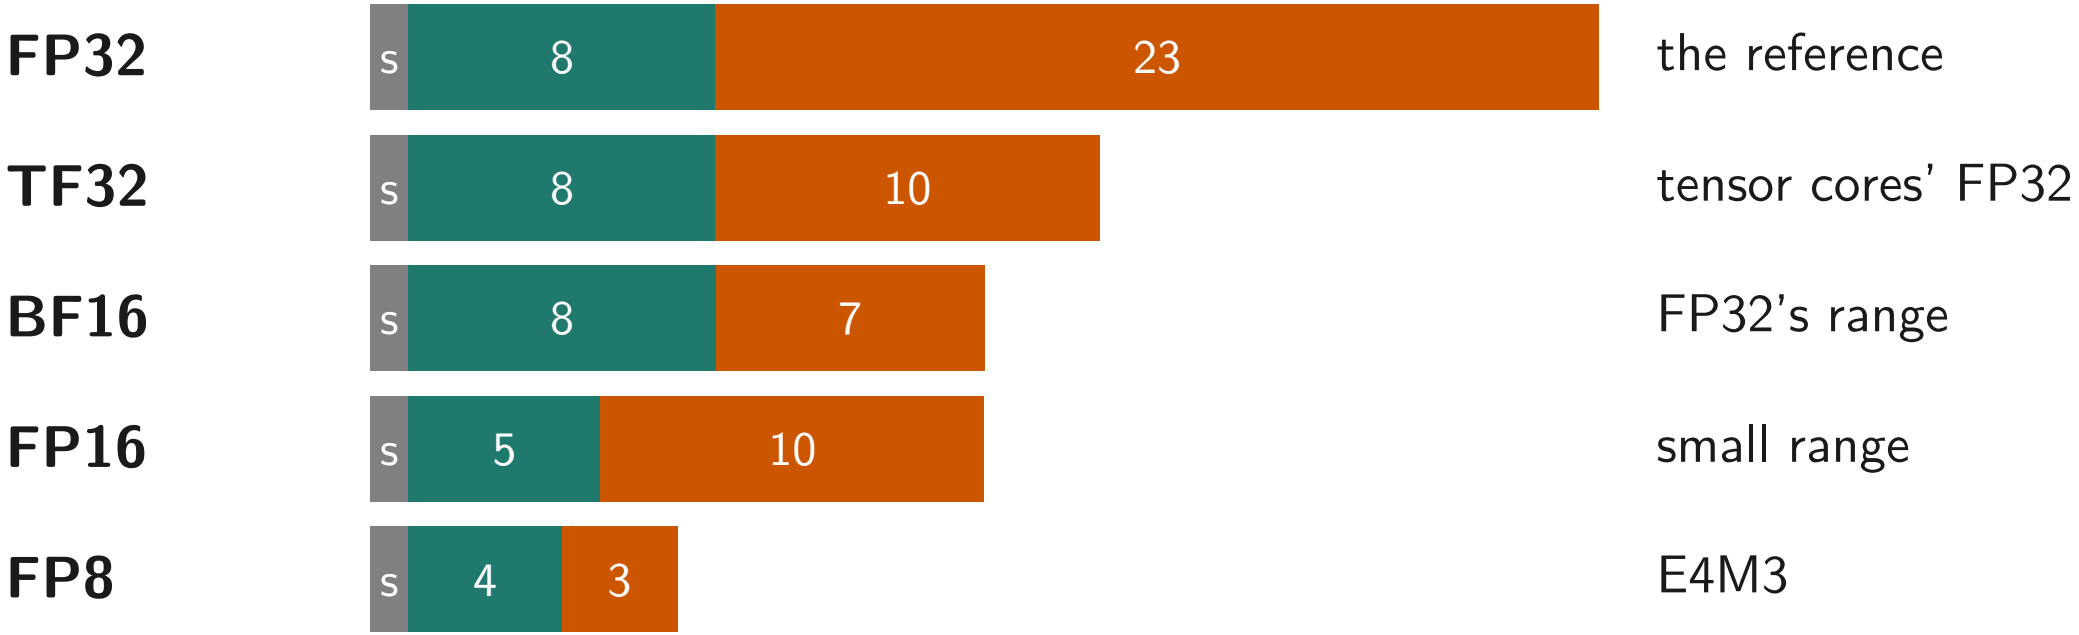



- **Exponent** bits $\rightarrow$ **range**: how large or small a number can be.
- **Mantissa** bits $\rightarrow$ **precision**: the gap after 1 is $\varepsilon = 2^{-23} \approx 1.2 \times 10^{-7}$.

Source: IEEE 754-2019 (FP32, FP16); NVIDIA A100 and H100 whitepapers (BF16, TF32); Micikevicius et al., 2022 (FP8).

# Five formats, bit by bit

grey: sign · teal: exponent bits · orange: mantissa bits · one cell = one bit




**Key idea**

BF16 keeps FP32's **range** and gives up **precision**; FP16 does the opposite.

Source: IEEE 754-2019 (FP32, FP16); NVIDIA A100 and H100 whitepapers (BF16, TF32); Micikevicius et al., 2022 (FP8).

# Range versus precision

| Format | Bits | Largest value | $\varepsilon$ (gap after 1) |
|---|---|---|---|
| FP32 | 32 | $\approx 3.4 \times 10^{38}$ | $1.2 \times 10^{-7}$ |
| BF16 | 16 | $\approx 3.4 \times 10^{38}$ | $7.8 \times 10^{-3}$ |
| FP16 | 16 | 65,504 | $9.8 \times 10^{-4}$ |
| FP8 (E4M3) | 8 | 448 | 0.125 |

**Key idea**

**Exponent bits buy range; mantissa bits buy precision.** With 16 bits you choose one; with 8 bits you need help (Part 2).

Source: IEEE 754-2019 (FP32, FP16); NVIDIA A100 and H100 whitepapers (BF16, TF32); Micikevicius et al., 2022 (FP8).

# Tensor cores and precision

| Format | Bits | H100 peak (dense) | Typical use |
|---|---|---|---|
| FP32 (ordinary lanes) | 32 | $\approx$ 67 TFLOP/s | Reference math |
| TF32 (tensor cores) | 19 | $\approx$ 495 TFLOP/s | Faster FP32 matmuls |
| BF16 / FP16 (tensor cores) | 16 | $\approx$ 989 TFLOP/s | Training default |
| FP8 (tensor cores) | 8 | $\approx$ 1,979 TFLOP/s | Large-model training |

**Watch out**

These peaks apply **only to matrix products**. Elementwise work (tanh, softmax, adds) runs on ordinary lanes and is usually memory-bound.

Source: H100 datasheet (SXM5); tensor-core values are the datasheet's sparse peaks divided by 2.

# Two ways low precision fails

**Underflow** (range)
gradient g $= 10^{-8}$
FP16: rounds to **0**
BF16, FP32: kept

**Swamping** (precision)
w + update $= 1 + 10^{-4}$
BF16, FP16: $=$ **1** (update lost)
FP32: kept

**Watch out**

Small numbers **vanish** (too little range) or small changes **vanish** (too little precision). Part 2 fixes both.

Source: IEEE 754-2019; checked in Python (NumPy float16, BF16 rounding emulated).

# Check yourself

Two minutes with a neighbour.

- An activation reaches 70,000. Which of FP16 and BF16 can store it? Why?
- What is the next FP8 (E4M3) number after 1? What does 1.1 become?
- A weight $w = 1$ is kept in BF16 and updated by $10^{-3}$. What happens?

Part 2

# Mixed precision

How do we train in 16 bits and still get FP32 answers?


1 · Number formats

**2** · Mixed precision

**3** · The anatomy of a matrix multiply

**Build** · Formats and tiles

# Multiply in BF16, add up in FP32

**BF16 inputs**
a[k], b[k]
2 bytes each

→

**product** a[k] · b[k]
exact

→

**FP32 accumulator**
s += a[k] · b[k]

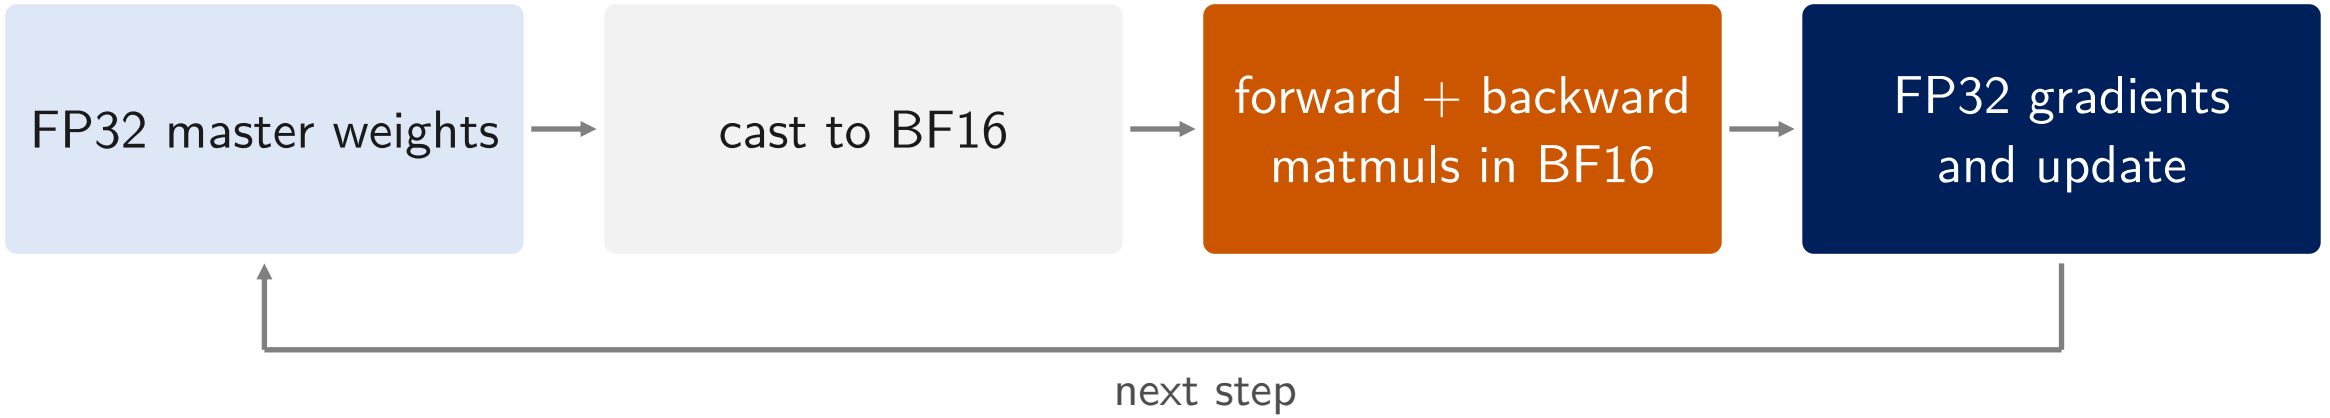
| Sum of 4,096 ones | Result |
|---|---|
| BF16 accumulator | **256** |
| FP32 accumulator | 4,096 |

**Key idea**

Round the **inputs**, never the **running sum**.

Source: NVIDIA PTX ISA, warp-level matrix multiply (mma.m16n8k16 for BF16, FP32 accumulator); H100 whitepaper.

# Mixed precision: the recipe




**Key idea**

Matrix multiplies in **BF16** for speed; sums, softmax, the loss, and the update in **FP32** for accuracy.

Source: Micikevicius et al., "Mixed Precision Training," ICLR 2018; PyTorch AMP docs.

# FP16 needs loss scaling; BF16 does not

$$\text{scaled loss: } \mathcal{S} \cdot L, \quad \text{for example } \mathcal{S} = 2^{16}$$

$$\text{chain rule: } \partial(\mathcal{S} \cdot L) \,/\, \partial w = \mathcal{S} \cdot \partial L \,/\, \partial w$$

$$10^{-8} \times 2^{16} \approx 6.6 \times 10^{-4} \quad \text{fits in FP16}$$

$$\text{before the update: divide by } \mathcal{S} \text{ (in FP32)}$$

**Check**

**BF16** has FP32's exponent range: gradients do not underflow, so it needs **no scaling**.

Source: PyTorch docs, "Automatic Mixed Precision" (torch.amp); Micikevicius et al., "Mixed Precision Training," ICLR 2018.

# FP8: a scale factor per tensor

$$\text{scale:} \quad s = 448 \;/\; \text{amax}(x)$$

$$\text{amax} = 7 \quad \Rightarrow \quad s = 64$$

$$\text{store} \;\; x \cdot s \;\; \text{in FP8;} \;\; \text{later divide by} \;\; s$$

**Why**

E4M3 holds about

**0.002 … 448**

Most tensors do not fit without a scale.

**Key idea**

FP8 = loss scaling **per tensor**: one FP32 scale each, so every tensor uses FP8's narrow range well.

Source: Micikevicius et al., "FP8 Formats for Deep Learning," 2022; NVIDIA Transformer Engine documentation.

# Mixed precision in PyTorch

```
model = MLP().cuda()              # FP32 weights
for x, y in loader:
    with torch.autocast("cuda",
                        dtype=torch.bfloat16):
        loss = F.cross_entropy(model(x), y)
    loss.backward()                  # FP32 .grad
    opt.step()                       # FP32 update
    opt.zero_grad()
```

- Wrap the **forward and loss** only.
- Parameters and gradients stay **FP32**.

**Key idea**

Autocast picks the format **per operation**: matrix multiplies go low, sums and the loss stay FP32.

Source: PyTorch docs, "Automatic Mixed Precision" (torch.amp); Micikevicius et al., "Mixed Precision Training," ICLR 2018.

# Check yourself

Two minutes with a neighbour.

- Why does a dot product of 4,096 BF16 products need an FP32 accumulator?
- Which format needs loss scaling, FP16 or BF16? Why?
- A tensor has amax $= 0.5$. What FP8 (E4M3) scale factor do you use?

Part 3

# The anatomy of a matrix multiply

Why does a naive matmul leave the GPU idle, and how do tiles fix it?

1 · Number formats | 2 · Mixed precision | **3** · The anatomy of a matrix multiply | **Build** · Formats and tiles

# Recall: FLOPs and the roofline

$$\mathcal{F} = 2 \cdot \mathcal{M} \cdot \mathcal{N} \cdot \mathcal{K} \qquad \text{(Week 11)}$$

$$\mathcal{T} \geq \max(\ \mathcal{F} \,/\, \pi\ ,\ \ \mathcal{Q} \,/\, \beta\ ), \qquad \mathcal{I} = \mathcal{F} \,/\, \mathcal{Q} \qquad \text{(Week 13)}$$

$$\text{ridge}\ \ \pi \,/\, \beta\text{:}\quad \text{FP32} \approx 20, \quad \text{BF16} \approx 295\ \ \text{FLOP/byte}$$

**Key idea**

A kernel is compute-bound only if it does more FLOPs per byte than the ridge. **What does a matmul kernel achieve?**

Source: H100 datasheet (SXM5); Williams, Waterman & Patterson, CACM 2009.

# A naive matmul is memory-bound

```
i, j = my_output()          # one C[i, j]
acc = 0.0
for k in range(K):
    acc += A[i, k] * B[k, j]    # 2 loads
C[i, j] = acc
```

$$\text{reads} \approx 2\text{MNK values}$$

$$I = 2\text{MNK} \;/\; (2\text{MNK} \cdot b) = 1 \;/\; b$$

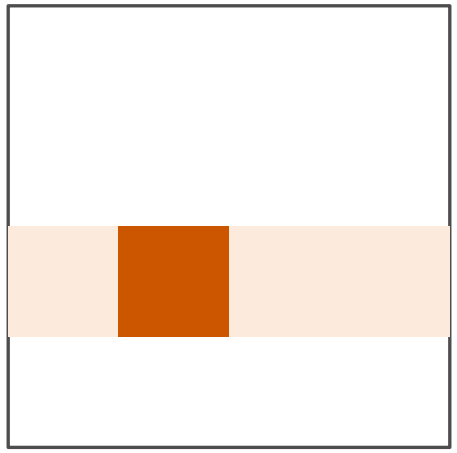
$$\text{FP32:}\quad I = 1/4 \;\ll\; 20$$

**Watch out**

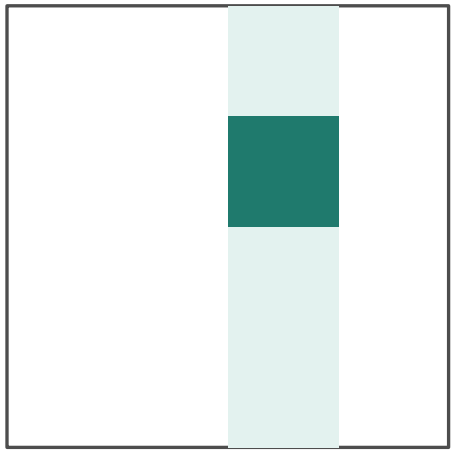
Every value is fetched from HBM **K times**. In FP32 the kernel reaches $\approx$ **0.84 TFLOP/s**: about **1%** of the FP32 peak.

Source: H100 datasheet (SXM5), 3.35 TB/s; derived values.

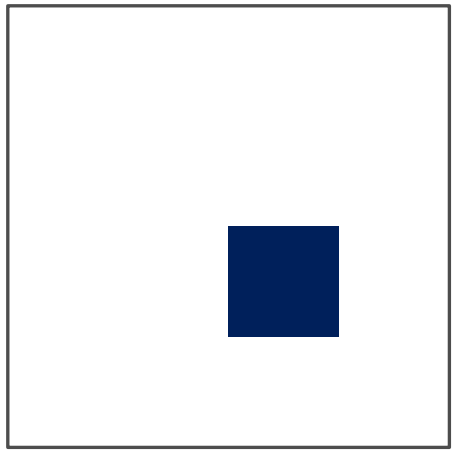
# Tiling: reuse data on chip



**A (M × K)**



**B (K × N)**



**C (M × N)**

**Naive**: each output re-reads a whole row of A and column of B from HBM.

**Tiled**: load one T × T tile of A and of B into shared memory; use each value **T times**.

**Key idea**

HBM traffic falls by a factor of ≈ **T**: tiles move once, then are reused on chip.

# Tiling, quantified

$$\text{naive:}\quad \text{reads} \approx 2\text{MNK values} \quad \Rightarrow \quad I = 1 \,/\, b$$

$$\text{tiled:}\quad \text{reads} \approx 2\text{MNK} \,/\, T\,\text{values} \quad \Rightarrow \quad I \approx T \,/\, b$$

$$\text{FP32},\ T = 64\text{:}\ \ I = 16 \qquad\qquad \text{BF16},\ T = 128\text{:}\ \ I = 64$$

**Watch out**

Lower precision raises the ridge to $\approx 295$: in BF16, tiling is **necessary**, and one level of it is **not enough**.

Source: Derived; ridge points from the H100 datasheet (SXM5).

# How big can a tile be?

$$\text{shared memory per } k\text{-step:} \quad (\mathrm{BM} \cdot \mathrm{BK} + \mathrm{BK} \cdot \mathrm{BN}) \cdot b$$

$$\mathrm{BM} = \mathrm{BN} = 128, \quad \mathrm{BK} = 64, \quad \mathrm{BF16}: \quad 32 \text{ KB}$$

$$\text{double-buffered:} \quad 64 \text{ KB} \quad \text{of } 228 \text{ KB per SM}$$

**Memory**

Tiles are capped by **shared memory and registers**. The **L2 cache** serves the tiles that neighbouring blocks re-read, so HBM sees less.

Source: NVIDIA H100 Tensor Core GPU datasheet (SXM5); NVIDIA H100 Tensor Core GPU Architecture whitepaper (2022).

# Register tiling: a second level




per $k$: $2r$ values in, $r^2$ multiply-adds

$$I = 2r^2 \,/\, (2r \cdot 4) = r \,/\, 4$$

$$r = 8: \quad 2 \times 33 \text{ TB/s} \approx 67 \text{ TFLOP/s}$$

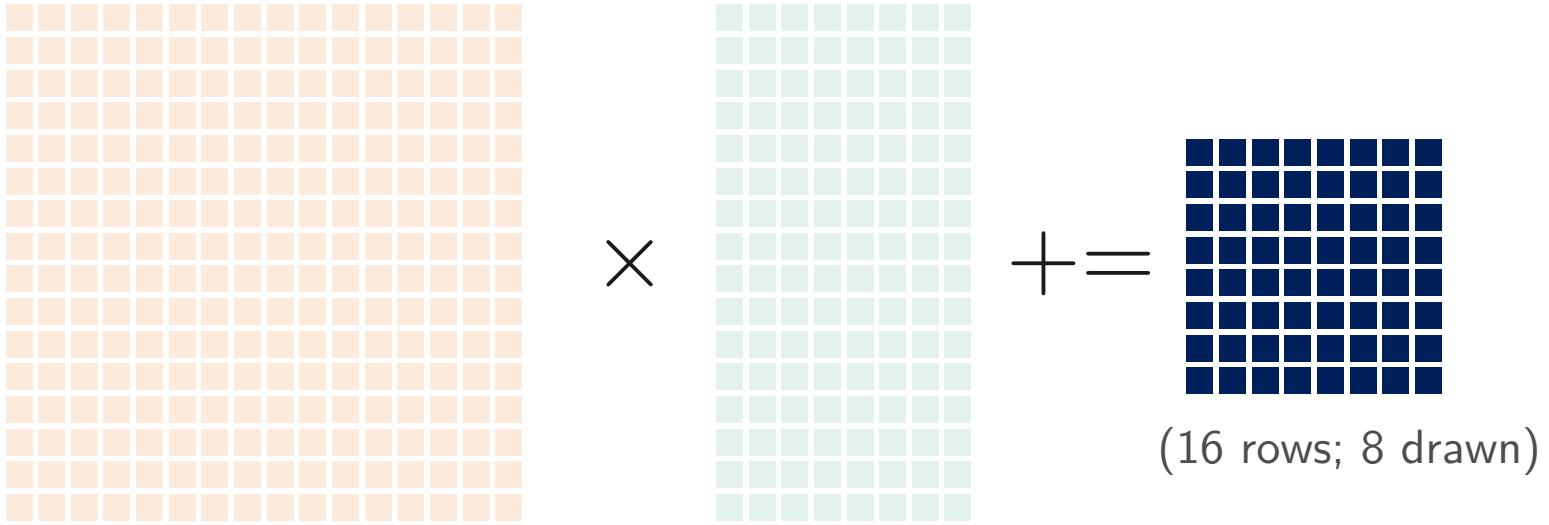
**Key idea**

Each level of memory needs its own reuse: **shared-memory tiles** for HBM, **register tiles** for shared memory.

Source: H100 whitepaper (128 bytes per cycle per SM, 132 SMs); derived, as in Week 13.

# Tensor cores multiply tiles

**A: 16 × 16 (BF16)** **B: 16 × 8 (BF16)** **C: 16 × 8 (FP32)**



(16 rows; 8 drawn)

**mma 16 × 8 × 16**

One instruction, one warp:

**4,096 FLOPs**

inputs BF16, sum **FP32**

**Key idea**

Tensor cores consume **small tiles from registers**. Shared-memory and register tiling exist to **feed** them.

Source: NVIDIA PTX ISA, warp-level matrix multiply (mma.m16n8k16 for BF16, FP32 accumulator); H100 whitepaper.

# Tile shapes: pad to a multiple

N = 322 columns, computed in 64-wide tiles



**384 computed; 62 are padding**

$$\text{64-wide tiles:} \quad \lceil 322 \,/\, 64 \rceil \cdot 64 = 384 \quad (16\% \text{ wasted})$$

$$\text{16-wide tiles:} \quad \lceil 322 \,/\, 16 \rceil \cdot 16 = 336 \quad (4\% \text{ wasted})$$

**Watch out**

Choose dimensions that are **multiples of the tile** (16, 64, 128): otherwise you pay for padding.

Source: NVIDIA Deep Learning Performance Guide, "Matrix Multiplication Background" (tile and wave quantization).

# Three levels of tiles, together

**HBM**
whole A, B, C

**Shared memory · per block**
$128 \times 64$ slices of A and B

**Registers · per warp**
$64 \times 64$ FP32 accumulators

**Tensor core · per instruction**
$16 \times 8 \times 16$, BF16 in, FP32 sum

smaller tiles,
faster memory,
more reuse

**Key idea**

Each level **reuses** what the level above delivered. Next lecture, Triton writes the lower levels for you.

# Check yourself

Two minutes with a neighbour.

- A naive BF16 matmul: what intensity? Memory- or compute-bound on an H100?
- What intensity do $T = 32$ tiles give in FP32? $T = 128$ tiles in BF16?
- Why make a layer's width 384 rather than 322?

IN-CLASS BUILD

# Formats and tiles

Work in pairs, on paper. About ten minutes.

1 · Number formats

2 · Mixed precision

3 · The anatomy of a matrix multiply

**Build** · Formats and tiles

# Exercise A: pick a format

| Part of the training step | Format? |
|---|---|
| Master weights and the update $w \leftarrow w - \eta \cdot g$ | |
| Inputs of the forward and backward matmuls | |
| Each matmul's running sum (accumulator) | |
| Softmax and the cross-entropy loss | |

**STRETCH**

The GPU has no BF16, only FP16. What must you add, and why?

# Exercise B: count a tiled matmul

$$\mathbf{C = A \cdot B},\ \ M = N = K = 4096,\ \ \text{BF16}$$
$$\text{tiles: } BM = BN = 128,\ \ BK = 64$$

- **1.** FLOPs F.
- **2.** Shared memory per block for one k-step of A and B.
- **3.** Bytes read into the SMs, and the intensity I.
- **4.** Memory- or compute-bound against the BF16 ridge? What T would reach it?

**USE**

$\beta = 3.35$ TB/s, $\pi = 989$ TFLOP/s, ridge $\approx 295$. Stretch: question 4 in FP32 (ridge $\approx 20$).

Source: NVIDIA H100 Tensor Core GPU datasheet (SXM5); NVIDIA H100 Tensor Core GPU Architecture whitepaper (2022).

# Solutions A: formats

| Part of the training step | Format |
|---|---|
| Master weights and the update | **FP32**: no swamping |
| Matmul inputs | **BF16**: tensor cores |
| Matmul accumulator | **FP32**: the puzzle |
| Softmax and loss | **FP32**: exp, log, sums |

**Check**

FP16 only: **loss scaling**, because its small range makes gradients underflow.

# Solutions B: a tiled matmul

| Question | Answer |
|---|---|
| 1. F | $2 \cdot 4096^3 \approx 1.37 \times 10^{11}$ FLOPs |
| 2. Shared memory | $(128 \cdot 64 + 64 \cdot 128) \times 2$ B $= 32$ KB (64 KB double-buffered) |
| 3. Bytes, I | $2MNK \cdot 2 / 128 \approx 2.1$ GB; $I = T / b = 64$ |
| 4. Bound? | $64 < 295$: memory-bound from HBM; needs $T \approx 590$ |

**Memory**

$T \approx 590$ would need a **1.4 MB** accumulator: impossible. The **L2 cache** closes the rest of the gap.

Source: NVIDIA H100 Tensor Core GPU datasheet (SXM5); NVIDIA H100 Tensor Core GPU Architecture whitepaper (2022).

# Key takeaways

- **Exponent** bits buy range, **mantissa** bits buy precision: BF16 keeps range, FP16 keeps precision.
- Mixed precision: **BF16 inputs**, **FP32 accumulation**, FP32 weights and update; FP16 adds **loss scaling**, FP8 a **scale per tensor**.
- A **naive** matmul has $\mathrm{I} = 1/\mathrm{b}$: **memory-bound**.
- **Tiling** gives $\mathrm{I} \approx \mathrm{T}/\mathrm{b}$; **register tiles** and **tensor cores** repeat the idea one level down.
- Pad dimensions to **multiples of the tile**.

# Next time, and readings

| When | What |
|---|---|
| After today | Lab 3 due |
| Lecture 2 (this week) | Triton, and one training step end to end |
| Friday, April 23 | T1 due |

**READ BEFORE LECTURE 2**

Micikevicius et al., "Mixed Precision Training," ICLR 2018.

NVIDIA, "Matrix Multiplication Background" (Deep Learning Performance Guide).

Tillet, Kung & Cox, "Triton," MAPL 2019 (for Lecture 2).